


# Kubernetes Administrator Bootcamp

*Become a Kube Astronaut*

*20 September 2025*



# Course Overview

1. Fundamentals & Architecture
2. Objects & Workloads
3. Networking & Services
4. Storage & Configuration
5. Resource Management & Health
6. Scheduling & Placement
7. Autoscaling & Performance
8. Security & Access Control
9. Extensions & Observability
10. Installation & Troubleshooting

# Why Kubernetes?

## Scalability

Manage hundreds of containers across nodes, scale up/down based on demand, and maintain service availability.

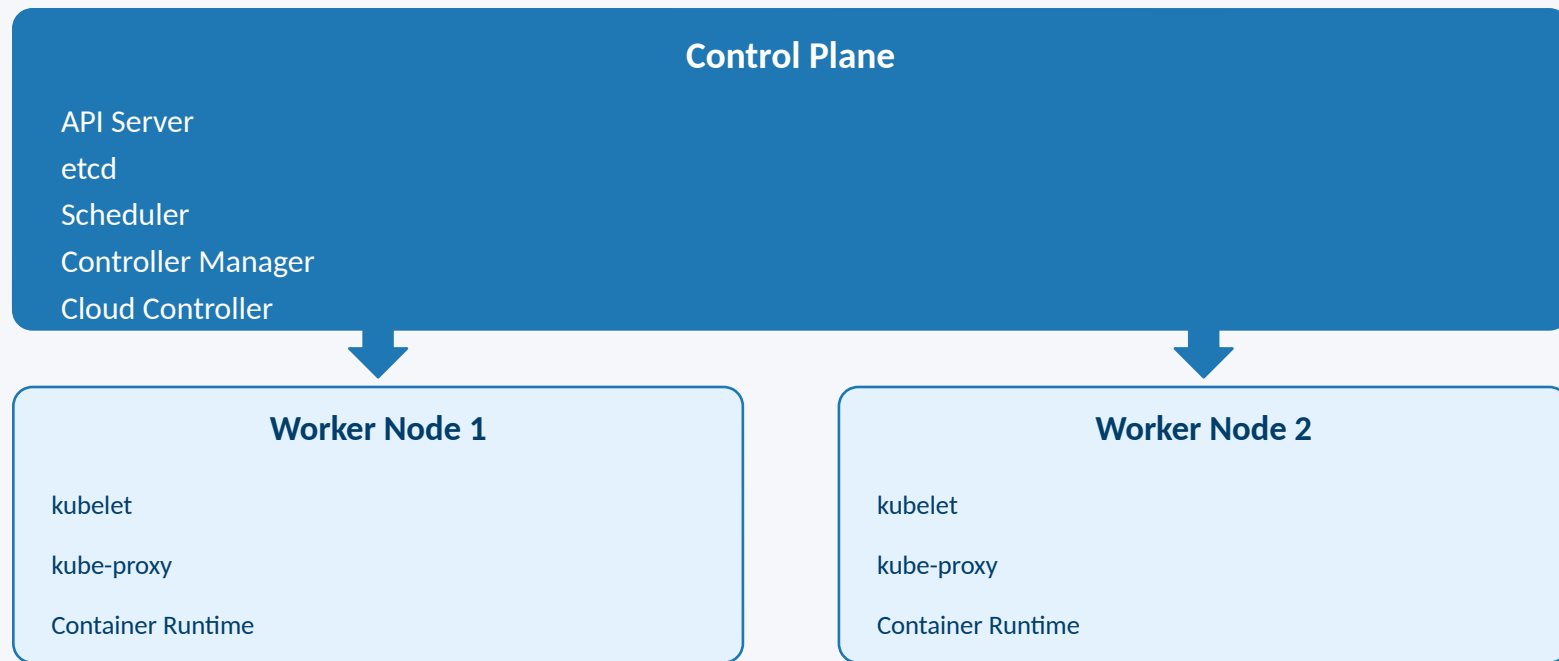
## Reliability

Self-healing and rolling updates keep applications running with minimal downtime and consistent releases.

## Portability

Runs on any infrastructure: on-prem, public cloud or hybrid. Declarative configs enable repeatable deployments.

# Cluster Architecture



# Control Plane Components

## API Server

Exposes the Kubernetes API. All control plane components and external clients interact through it.

## Scheduler

Watches for new Pods and assigns them to nodes based on resource requirements, constraints and affinity rules.

## Cloud Controller

Integrates Kubernetes with cloud-provider APIs for load balancers, persistent volumes and node lifecycle.

## etcd

Highly-available key-value store backing all cluster data and configuration.

## Controller Manager

Runs controllers that handle routine tasks such as node health, job management and endpoint reconciliation.

# Node Components

## **kubelet**

Agent running on each node that ensures containers described in Pod specs are running and healthy.

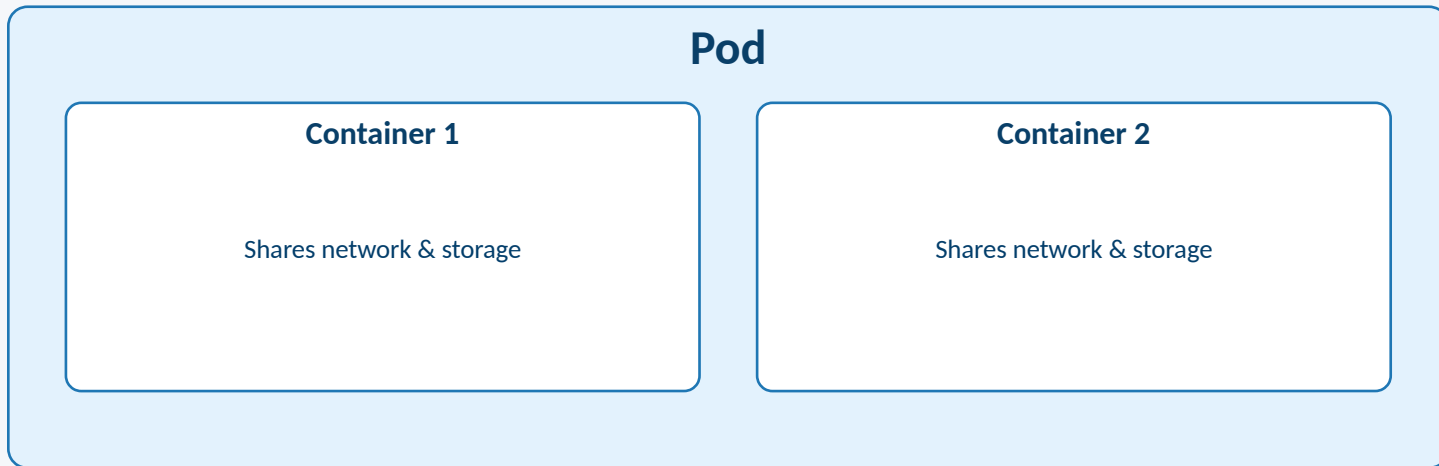
## **kube-proxy**

Implements the Service concept by managing network rules on each node and proxying requests to Pods.

## **Container Runtime**

Software such as containerd or CRI-O that runs and manages containers as instructed by the kubelet.

# Pods: Smallest Deployable Unit



**Pods** are the smallest deployable unit in Kubernetes. A pod represents one or more containers that share the same IP address, storage volumes and namespaces.

# ReplicaSets & Deployments

## ReplicaSet

Maintains a specified number of replica Pods

- 
- Replaces Pods if they crash or are deleted
- 
- Used by Deployments

## Deployment

- Manages ReplicaSets and rollout strategies
- Performs rolling updates: at least 75% pods available, up to 125% pods during updates
- Enables rollbacks and versioning



# StatefulSet, DaemonSet & Jobs

## StatefulSet

Manages stateful applications with stable network identities and persistent volumes.

## DaemonSet

Ensures a copy of a Pod runs on every node (or selected nodes). Useful for logging, monitoring agents.

## Jobs/CronJobs

Runs one-off or scheduled tasks. Jobs ensure completion; CronJobs schedule repeating runs.

# Services & Networking

## ClusterIP

Exposes a service on a cluster-internal IP. Reachable only within the cluster (default).

## NodePort

Exposes the service on each node's IP and static port. Useful for simple external access.

## LoadBalancer

Creates an external load balancer in cloud environments and assigns a public IP.

## ExternalName

Maps the service to an external DNS name via a CNAME without proxying.

# Ingress & HTTP Routing

**Ingress** is an API object that exposes HTTP and HTTPS routes from outside the cluster to Services within the cluster. It may provide load balancing, SSL/TLS termination and name-based virtual hosting. An Ingress controller fulfills these rules. Only creating an Ingress resource has no effect without a controller.

## Ingress Rule


Host:  
foo.example.com

Path:  
/test → Service A  
  
/api → Service B

# Persistent Volumes & Claims

## PersistentVolume (PV)

- Cluster-level resource representing storage provisioned by administrators or dynamic provisioners.
- Lifecycle independent of Pods.
- PersistentVolumeClaim (PVC)
- User request for storage specifying size and access mode.
- Binds to a matching PV and consumed by a Pod like a node resource.



# ConfigMaps & Secrets

## ConfigMap

- Stores non-confidential key-value pairs.
- Decouples environment configuration from container images.
- Consumed via environment variables, command args or mounted files.

## Secret

- Stores sensitive data such as passwords, tokens or SSH keys.
- Reduces the risk of accidentally exposing confidential data in images or manifests.
- Similar usage to ConfigMaps but contents should be protected (RBAC, encryption at rest).


# Resource Requests & Limits

## Requests

- Minimum amount of CPU/Memory reserved by the scheduler for the Pod.
- Determines which node the Pod can be scheduled on.

## Limits

- Maximum CPU/Memory usage enforced by the kubelet using cgroups.
- Memory limits trigger OOM kill on pressure; CPU limits throttle usage.
- If not specified, limit = request.



# Pod Quality of Service (QoS)

## Guaranteed

Every container has CPU & memory requests equal to limits; highest priority and reserved resources; cannot exceed limits.

## Burstable

At least one container has a request but limits may be higher; gets guaranteed minimum resources but may use excess; intermediate eviction priority.

## BestEffort

No requests or limits; uses resources only when available; lowest priority and first to be evicted under pressure.

# Probes: Liveness, Readiness & Startup

## Liveness

Detects if the application has entered a deadlock or crashed.

- 
- When failed repeatedly, kubelet restarts the container.
- 

## Readiness

Determines when a container is ready to accept traffic.

- 
- If fails, pod is temporarily removed from service endpoints.
- 

## Startup

Verifies the application has started.

- 
- Disables liveness/readiness checks until it succeeds, avoiding premature restarts.
-



# Scheduling: Labels & Affinity

## nodeSelector

- Simplest way to constrain a pod to run only on nodes with matching labels.
- Example: `nodeSelector: { disktype: ssd }`

## Node Affinity

- Offers more expressive placement constraints using required or preferred rules.
- Example: `requiredDuringSchedulingIgnoredDuringExecution` with label selector expressions.


# Taints & Tolerations

## Taints

- Applied to nodes to repel pods.
- Consist of key, value and effect (NoSchedule, PreferNoSchedule, NoExecute).


## Tolerations

- Pods specify tolerations to match taints on nodes, allowing scheduling or continued running.
- Effect defines behaviour: NoSchedule prevents scheduling; PreferNoSchedule soft repels; NoExecute evicts existing pods.



# Horizontal Pod Autoscaler (HPA)

- Automatically scales the number of pod replicas based on observed metrics (CPU, memory, custom metrics).
- Runs a control loop that periodically queries the metrics API and adjusts replicas accordingly.
- Works with Deployments, StatefulSets and other scalable resources.



# Vertical & Node Autoscaling

## Vertical Pod Autoscaler (VPA)


- Separate project that recommends and automatically updates CPU/Memory requests and limits for Pods.
- Modes: Auto (apply), Recreate (restart to apply), Initial (set once), Off.

## Node Autoscaling

- Cluster Autoscaler adds/removes nodes to meet pending Pod requirements.
- Karpenter auto-provisions nodes with just-in-time capacity and works across cloud providers.
- Works together with HPA to balance cost and capacity.

# Role-Based Access Control (RBAC)

- Role: set of permissions within a namespace.
- ClusterRole: permissions across the entire cluster.
- RoleBinding: grants Role to users/groups/service accounts within a namespace.
- ClusterRoleBinding: binds a ClusterRole cluster-wide.



# Service Accounts

- Non-human identities used by Pods and controllers to authenticate to the API server and external services.
- Namespaced and lightweight objects; each namespace has a default service account assigned to new Pods.
- Service accounts are distinct from user accounts and use tokens mounted into Pods for authentication.
- Use RBAC to assign least-privilege roles to service accounts.

# Pod Security Standards & Admission

<b>Privileged</b>	Unrestricted policy allowing widest possible permissions; intended for trusted workloads.
<b>Baseline</b>	Minimally restrictive; prevents known privilege escalations while allowing common configurations.
<b>Restricted</b>	Heavily restricted; follows current best practices for pod hardening.

## Pod Security Admission Modes


- enforce: reject pods that violate the selected policy.
- audit: allow the pod but record violation annotation in the audit log.
- warn: allow the pod but return a warning to

the user.**Namespace labels**`pod-security.kubernetes.io/<MODE>: <LEVEL>`

Example: `pod-security.kubernetes.io/enforce: restricted`

# Helm - Kubernetes Package Manager

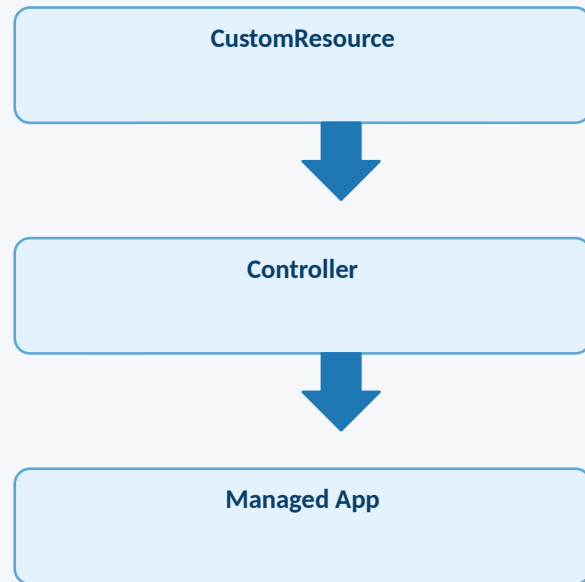
- Helm uses Charts to package all Kubernetes resources needed to deploy an application.
- Simplifies installation, upgrades and rollback of complex apps; handles dependencies and values templating.
- Supports chart repositories and sharing of community charts.





# Operators: Extend Kubernetes

- Operators are software extensions that use Custom Resources (CRs) to manage applications and their components.
- They follow Kubernetes principles: controllers watching desired state vs. actual state (control loop).
- Examples: automate application deployment, backup/restore, upgrades and leader election.




# Observability with Prometheus & Grafana

## Prometheus

- Open-source systems monitoring and alerting toolkit originally built at SoundCloud.
- Collects metrics as time-series data, identified by metric names and labels.
- Pulls metrics from targets via HTTP and stores them locally; provides PromQL for flexible queries.

## Grafana

- Open-source analytics and visualization platform.
- Connects to multiple data sources including Prometheus, InfluxDB and Elasticsearch.
- Enables interactive dashboards, alerts and plug-in extensibility.



# Debugging & Troubleshooting

## kubectl describe

Fetches detailed information about Pods, including status, resource requests/limits, events and conditions. Useful for diagnosing scheduling and lifecycle issues.

## kubectl logs

Prints stdout/stderr from containers. Use --previous to view logs from crashed containers.

## kubectl exec

Runs commands inside a running container for investigation. For example:  
kubectl exec -it <pod> -- sh.

## kubectl debug

Creates an Ephemeral container or copies a Pod with modified settings when your container lacks debugging utilities.

## Troubleshooting Workflow

1. Check Pod status and events with kubectl get & describe.
2. Inspect container logs. Use --previous for crash loops.
3. Exec into the container for manual inspection.
4. Create debug Pods when images lack troubleshooting tools.

# Cluster Installation with kubeadm

1. Prepare machines (install container runtime, kubeadm, kubelet, kubectl).
2. Initialize control plane: `kubect init ...` – downloads and installs control plane components and performs pre-checks.
3. Save kubeadm join command and configure kubeconfig (admin.conf) for kubectl access.
4. Deploy a CNI plugin: `kubectl apply -f <podnetwork>.yaml`.
5. Join worker nodes: `kubeadm join <control-plane-host>:<port> --token <token> --discovery-token-ca-cert-hash <hash>`.

## Notes & Cautions

- kubeadm init runs prechecks and generates admin.conf and super-admin.conf; do not share these files.
- Always record the join command and secure the bootstrap token.
- Install only one CNI plugin per cluster; ensure Pod network CIDR does not overlap with host networks.

# High Availability & Upgrades

## High Availability

- Run multiple control-plane nodes behind a load balancer.
- Use stacked etcd cluster for durability (or external managed etcd).
- Ensure quorum and monitor etcd health.

## Upgrades & Maintenance

- Use kubectl upgrade to incrementally upgrade control plane and nodes.
- Drain nodes before upgrading to move Pods to other nodes (kubectl drain).
- Back up etcd regularly and test restore procedures.

# Mission Accomplished

You are now equipped to administer Kubernetes clusters!

- Understand cluster architecture and core components.
- Deploy and scale workloads using Pods, Deployments, and Services.
- Manage storage, configuration and secrets effectively.
- Tune performance through resource management and autoscaling.
- Secure your cluster with RBAC, Pod Security and admission controls.
- Extend Kubernetes with Helm and Operators and monitor with Prometheus & Grafana.
- Install and troubleshoot clusters with kubeadm.

## Next Steps:

Explore advanced networking and security features, contribute to the Kubernetes community, and consider pursuing Kubernetes Administrator (CKA) certification.

[\[43\]](#)[\[46\]](#)[\[47\]](#)[\[48\]](#)[\[49\]](#)[\[50\]](#)[\[51\]](#)[\[52\]](#)[\[53\]](#)[\[54\]](#)

[\[55\]](#)[\[56\]](#)[\[57\]](#)[\[58\]](#)[\[59\]](#)



# Imperative vs Declarative

## Imperative

- `kubectl create/delete` operations act directly on cluster objects
- `kubectl run nginx --image=nginx --restart=Never`
- `kubectl create deployment web --image=nginx`

## Declarative

- Define desired state in YAML manifests (Pod, Deployment, Service)
- `kubectl apply -f pod.yaml`
- `kubectl apply -f deployment.yaml`
- `kubectl apply -f service.yaml`

# Exercises: Cluster & Nodes

- Start a local cluster: minikube start or kind create cluster
- List nodes: kubectl get nodes
- Inspect node details: kubectl describe node <node-name>
- Label a node: kubectl label node <node-name> disktype=ssd
- Taint a node: kubectl taint node <node-name> key=value:NoSchedule



# Exercises: Pods & Deployments

## Pods

- `kubectl run nginx --image=nginx --restart=Never`
- `kubectl get pods`
- Create `pod.yaml` with `apiVersion`, `kind`, `metadata` & `spec`
- `kubectl apply -f pod.yaml`

## Deployments

- `kubectl create deployment web --image=nginx --replicas=3`
- `kubectl scale deployment/web --replicas=5`
- `kubectl set image deployment/web nginx=nginx:1.17`
- Define `deployment.yaml` and apply: `kubectl apply -f deployment.yaml`
- `kubectl rollout history deployment/web`
- `kubectl rollout undo deployment/web`

# Exercises: Services & Ingress

## Services

- Expose deployment: `kubectl expose deployment web --port=80 --target-port=80 --type=ClusterIP`
- NodePort service: `kubectl expose deployment web --port=80 --target-port=80 --type=NodePort`
- List services: `kubectl get services`
- Define service.yaml and apply: `kubectl apply -f service.yaml`

## Ingress

- Create ingress.yaml with host and path rules
- Apply: `kubectl apply -f ingress.yaml`
- Check ingress: `kubectl get ingress`
- Describe ingress: `kubectl describe ingress <name>`

# Exercises: Storage & Configuration

## ConfigMaps & Secrets

- `kubectl create configmap app-config --from-literal=APP_MODE=production`
- `kubectl create secret generic db-secret --from-literal=password=Pa$$w0rd`
- `kubectl get configmap,secret`
- `kubectl describe configmap app-config`
- Define configmap.yaml/secret.yaml & apply: `kubectl apply -f configmap.yaml`

## Persistent Storage

- Define pv.yaml for static PersistentVolume
- Define pvc.yaml for PersistentVolumeClaim
- Apply PV & PVC: `kubectl apply -f pv.yaml && kubectl apply -f pvc.yaml`
- View storage: `kubectl get pv,pvc`

# Exercises: Resource Management & Health

## Requests & Limits

- Set resources in pod spec (requests & limits) and apply via YAML
- `kubectl set resources deployment/web --limits=cpu=500m,memory=256Mi --requests=cpu=250m,memory=128Mi`
- `kubectl describe pod <pod>` (view QoS class & resource usage)
- `kubectl top pods` (requires metrics-server)

## Liveness/Readiness Probes

- Add `livenessProbe/readinessProbe/startupProbe` in container spec
- `kubectl apply -f deployment-probes.yaml`
- `kubectl describe pod <pod>` to view probe status

# Exercises: Scheduling & Placement

## Labels & Node Selectors

- Label nodes: `kubectl label nodes <node> disktype=ssd zone=east`
- Add nodeSelector to Pod: `nodeSelector: { disktype: ssd }`
- `kubectl apply -f pod-selector.yaml`

## Affinity & Taints

- Define `nodeAffinity/podAffinity/podAntiAffinity` in Pod spec
- `kubectl taint nodes <node> dedicated=database:NoSchedule`
- Add tolerations: `tolerations: [{ key: 'dedicated', operator: 'Equal', value: 'database', effect: 'NoSchedule' }]`
- Apply Pod with tolerations: `kubectl apply -f pod-tolerations.yaml`

# Exercises: Autoscaling & Performance

## Horizontal Pod Autoscaler

- `kubectl autoscale deployment web --min=2 --max=5 --cpu-percent=50`
- `kubectl get hpa`
- Define `hpa.yaml` with target metrics and apply

## Load Generation & VPA

- Generate load: `kubectl run -i --tty load-gen --image=busybox -- /bin/sh`
- Use stress tools inside the container to increase CPU usage
- Delete load generator: `kubectl delete pod load-gen`
- Install VPA components and apply `vpa.yaml` (optional)
- Use cluster autoscaler/Karpenter for node scaling (cloud-provider specific)

# Exercises: Security & RBAC

## RBAC

- Create service account: `kubectl create serviceaccount viewer`
- Define Role manifest (read pods) and apply
- Bind role: `kubectl create rolebinding read-pods-binding --role=pod-reader --serviceaccount=default:viewer --namespace=default`
- Test access: `kubectl auth can-i list pods --as system:serviceaccount:default:viewer`

## Pod Security

- Label namespace to enforce baseline: `kubectl label --overwrite ns dev pod-security.kubernetes.io/enforce=baseline`
- Audit with restricted profile: `kubectl label --overwrite ns dev pod-security.kubernetes.io/audit=restricted`
- Warn on violations: `kubectl label --overwrite ns dev pod-security.kubernetes.io/warn=baseline`
- Verify labels: `kubectl describe ns dev`

# Exercises: Extensions & Observability

## Helm & Operators

- `helm repo add bitnami https://charts.bitnami.com/bitnami`
- `helm install my-nginx bitnami/nginx`
- `helm upgrade my-nginx bitnami/nginx --set service.type=NodePort`
- `helm uninstall my-nginx`
- Install operator: `kubectl apply -f operator.yaml`
- List custom resources: `kubectl get crd`
- Apply custom resource instance: `kubectl apply -f cr-instance.yaml`

## Observability

- `helm repo add prometheus-community https://prometheus-community.github.io/helm-charts`
- `helm install kube-prom prometheus-community/kube-prometheus-stack`
- `kubectl get pods -n kube-prometheus-stack`
- `kubectl port-forward svc/kube-prometheus-stack-grafana -n kube-prometheus-stack 3000:80`
- View logs: `kubectl logs <pod>`
- Follow logs: `kubectl logs -f <pod>`



# Exercises: Installation & Troubleshooting

## Cluster Installation

- `kubeadm init --pod-network-cidr=192.168.0.0/16`
- `mkdir -p $HOME/.kube && cp /etc/kubernetes/admin.conf $HOME/.kube/config`
- `kubectl get nodes`
- `kubeadm join <master-ip>:6443 --token <token> --discovery-token-ca-cert-hash sha256:<hash>`

## Maintenance & Troubleshooting

- cordon & drain node: `kubectl cordon <node>`, `kubectl drain <node>`
- Upgrade plan/apply: `kubeadm upgrade plan`; `kubeadm upgrade apply`
- Describe resources: `kubectl describe pod <pod>`
- View logs: `kubectl logs <pod> [--previous]`
- Execute into container: `kubectl exec -it <pod> -- /bin/sh`
- Debug pod with temporary container: `kubectl debug <pod> --image=busybox`